

DSCI 6607– Fall 2024 Assignment 4

2024-10-26

```
# Abdallah Chidjou
# Citation: (source of help: Lecture note, googling in general, stackoverflow, and chatgpt)
```

Question 1

```
# Load mtcars (this is already available in R)
data(mtcars)
# View the first few rows of the dataset
head(mtcars)
```

```
##           mpg  cyl  disp  hp  drat    wt   qsec vs  am  gear  carb
## Mazda RX4      21.0   6  160 110 3.90 2.620 16.46 0   1    4    4
## Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02 0   1    4    4
## Datsun 710      22.8   4  108  93 3.85 2.320 18.61 1   1    4    1
## Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44 1   0    3    1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02 0   0    3    2
## Valiant        18.1   6  225 105 2.76 3.460 20.22 1   0    3    1
```

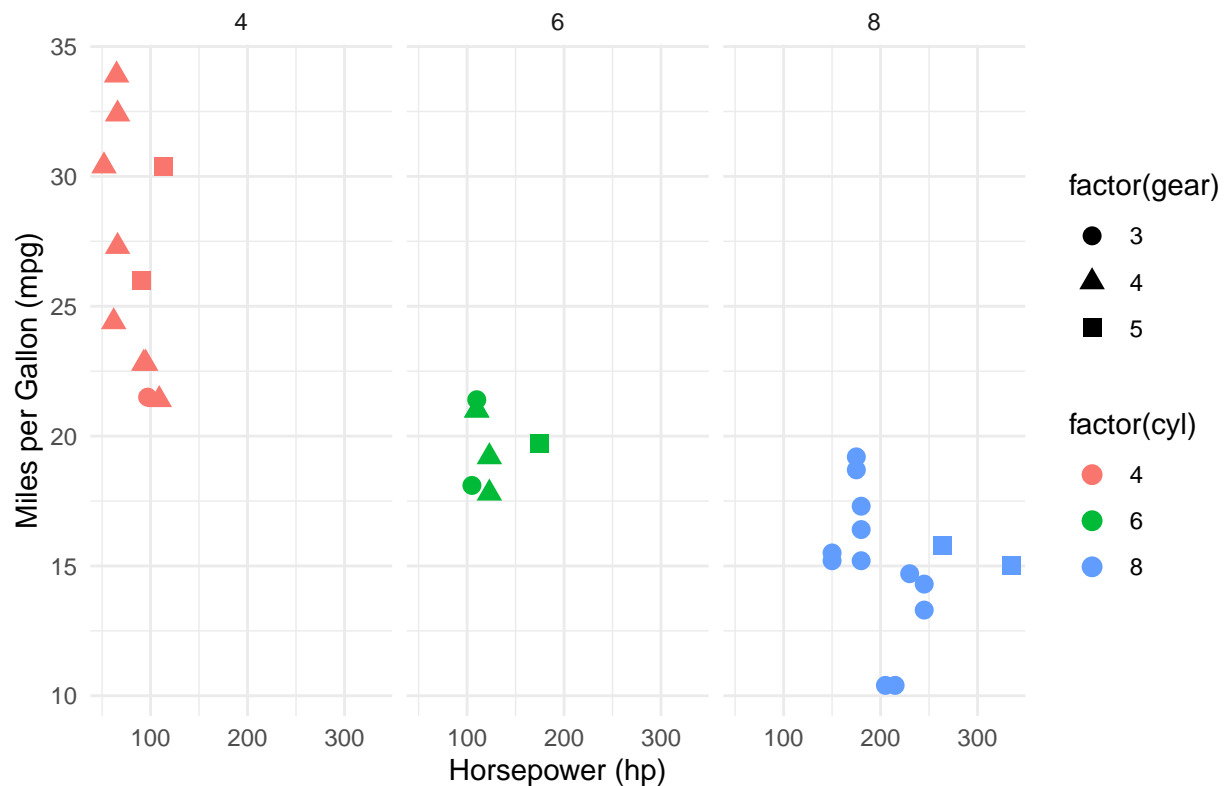
```
# View the structure of the dataset to understand the variable types
#str(mtcars)
```

```
# Get a summary of the dataset
#summary(mtcars)
```

```
# Load the necessary library
library(ggplot2)
```

```
# Create the scatter plot
ggplot(mtcars, aes(x = hp, y = mpg)) +
  geom_point(aes(color = factor(cyl), shape = factor(gear)), size = 3) +
  facet_wrap(~ cyl) +
  theme_minimal() +
  labs(title = "Scatter Plot of MPG vs Horsepower, Faceted by Cylinders",
       x = "Horsepower (hp)",
       y = "Miles per Gallon (mpg)")
```

Scatter Plot of MPG vs Horsepower, Faceted by Cylinders



```
# a.
# Each panel represents a subset of cars from the dataset
# with the same number of cylinders (cyl = 4, 6, or 8).

# b.
# geom_point(aes(color = factor(cyl), shape = factor(gear)), size = 3)
# done in the plot

# c.
# The variable used for faceting is cyl.

# d.
# Cars with more cylinders (e.g., 8 cylinders) tend to have lower mpg and higher hp,
# indicating they are powerful but less efficient.
# Cars with fewer cylinders (e.g., 4 cylinders) tend to have higher mpg and lower hp,
# indicating they are more fuel-efficient but less powerful.
```

Question 2

```
# a1: Parsing and outputting the date
a1 <- "12/30/14"
parsed_a1 <- strptime(a1, format = "%m/%d/%y")
formatted_a1 <- format(parsed_a1, format = "%b %d, %Y")
print(formatted_a1)
```

```
## [1] "Dec 30, 2014"
```

```
# a2: Parsing and outputting the date
a2 <- "07-Jan-2017"
parsed_a2 <- strptime(a2, format = "%d-%b-%Y")
formatted_a2 <- format(parsed_a2, format = "%d-%m-%Y")
print(formatted_a2)
```

```
## [1] "07-01-2017"
```

```
# Create the original vector of date-time strings
convert_date_time <- function(date_vector) {
  # Use sapply to iterate over the vector and parse/format each element
  formatted_dates <- sapply(date_vector, function(x) {
    # Parse the original string to a date-time object
    parsed_date <- strptime(x, format = "%B %d (%Y) - %I:%M%p")
    # Format the parsed date into the desired format
    formatted_date <- format(parsed_date, format = "%m/%d/%Y - %I:%M%p")
    return(formatted_date)
  })
  # Return the formatted dates
  return(formatted_dates)
}

# a3: Parsing and outputting the date
a3 <- c("August 19 (2015) - 3:04PM", "July 1 (2015) - 4:04PM")
formatted_dates <- convert_date_time(a3)
# Print the formatted dates
print(formatted_dates)
```

```
## August 19 (2015) - 3:04PM      July 1 (2015) - 4:04PM
##      "08/19/2015 - 03:04PM"    "07/01/2015 - 04:04PM"
```

```
# a4: Parsing and outputting the date
a4 <- "January 1, 2010"
parsed_a4 <- strptime(a4, format = "%B %d, %Y")
print(parsed_a4)
```

```
## [1] "2010-01-01 NST"
```

```
# a5: Parsing and outputting the date
a5 <- "2015-Mar-07"
parsed_a5 <- strptime(a5, format = "%Y-%b-%d")
print(parsed_a5)
```

```
## [1] "2015-03-07 NST"
```

Question 3

```
# Install and load the required libraries
# install.packages(c("dplyr", "ggrepel"))
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(ggplot2)
library(ggrepel)
```

```
# Load the dataset (assuming you have the CSV file named "largest_cities.csv")
cities <- read.csv("largest_cities.csv")
```

```
# a. Create a new column `city_density`
cities <- cities %>%
  mutate(city_density = city_pop / city_area)
head(cities)
```

```
##      name country      city_definition population city_pop city_area
## 1   Tokyo  Japan  Metropolis prefecture    37.400   13.515    2191
## 2    Delhi India National capital territory    28.514   16.753    1484
## 3 Shanghai China      Municipality        25.582   24.183    6341
## 4  São Paulo Brazil      Municipality        21.650   12.252    1521
## 5 Mexico City Mexico      City-state        21.581    8.919    1485
## 6   Cairo  Egypt      Urban governorate     20.076    9.500    3085
## metro_pop metro_area urban_pop urban_area      wiki country_code2
## 1   37.274    13452    38.505     8223    Tokyo          JP
## 2   29.000     3483    28.125     2240    Delhi           IN
## 3      NA       NA    22.125     4015    Shanghai        CN
## 4   21.735     7947    20.935     3043    Sao_Paulo        BR
## 5   20.893     7854    20.395      237 Mexico_City       MX
## 6      NA       NA    16.925     1917    Cairo           EG
## country_code3      country_name_official      continent      lon
## 1          JPN              Japan          Asia 139.69222
## 2          IND      India, Republic of      Asia  77.23000
## 3          CHN  China, People's Republic of      Asia 121.47472
## 4          BRA Brazil, Federative Republic of South America -46.63333
## 5          MEX Mexico, United Mexican States North America -99.13333
## 6          EGY      Egypt, Arab Republic of      Africa 31.22889
##      lat koppen_code koppen_main      city num cost_of_living
## 1  35.68972      Cfa  Temperate    Tokyo, Japan   18      86.87
## 2  28.61000      BSh      Dry    Delhi, India  405      28.18
```

```
## 3 31.22861 Cfa Temperate Shanghai, China 235 50.07
## 4 -23.55000 Cfb Temperate Sao Paulo, Brazil 257 45.52
## 5 19.43333 Cwb Temperate Mexico City, Mexico 317 38.55
## 6 30.05806 BWh Dry Cairo, Egypt 393 30.94
## cost_rent cost_groceries cost_restaurant local_pp city_density
## 1 38.00 83.42 56.70 89.70 0.006168416
## 2 8.18 26.15 24.76 54.69 0.011289084
## 3 35.67 52.50 36.48 54.40 0.003813752
## 4 16.28 33.29 39.70 31.97 0.008055227
## 5 20.43 33.75 33.11 42.91 0.006006061
## 6 6.48 26.41 25.43 24.53 0.003079417
```

```
# b. Select name, city_pop, city_area, and city_density columns
selected_data <- cities %>%
  select(name, city_pop, city_area, city_density)

# Display the selected data
# print(selected_data)

# c. Modify city_density by multiplying by 1000
cities <- cities %>%
  mutate(city_density = city_density * 1000)

# Update the `selected_data` after modifying city_density
selected_data <- cities %>%
  select(name, city_pop, city_area, city_density)

# Display the updated selected data
print(selected_data)
```

```
##          name city_pop city_area city_density
## 1      Tokyo  13.515   2191.00   6.16841625
## 2      Delhi  16.753   1484.00  11.28908356
## 3    Shanghai  24.183   6341.00   3.81375177
## 4    São Paulo  12.252   1521.00   8.05522682
## 5 Mexico City   8.919   1485.00   6.00606061
## 6      Cairo   9.500   3085.00   3.07941653
## 7      Mumbai  12.478    603.00  20.69320066
## 8      Beijing  21.707  16411.00   1.32271038
## 9        Dhaka  14.399    338.00  42.60059172
## 10       Osaka   2.725    225.00  12.11111111
## 11 New York City   8.399    786.00  10.68575064
## 12      Karachi  14.910    378.00  39.44444444
## 13 Buenos Aires   3.054    203.00  15.04433498
## 14    Chongqing  30.166  82403.00   0.36607890
## 15     Istanbul  15.029   5196.00   2.89241724
## 16     Kolkata   4.497    205.00  21.93658537
## 17 Metro_Manila   1.780     43.00  41.39534884
## 18       Lagos    NA        NA        NA
## 19 Rio de Janeiro   6.520   1221.00   5.33988534
## 20      Tianjin  15.569   1192.00  13.06124161
## 21    Kinshasa  11.462   9965.00   1.15022579
## 22    Guangzhou  14.498   7434.00   1.95022868
## 23   Los Angeles   3.990   1214.00   3.28665568
```

## 24	Moscow	13.200	2511.00	5.25686977
## 25	Shenzhen	12.528	205.00	61.11219512
## 26	Lahore	11.126	1772.00	6.27878104
## 27	Bangalore	8.444	709.00	11.90973202
## 28	Paris	2.148	105.00	20.45714286
## 29	Bogotá	7.963	1587.00	5.01764335
## 30	Jakarta	10.154	664.00	15.29216867
## 31	Chennai	6.727	426.00	15.79107981
## 32	Lima	8.894	2672.00	3.32859281
## 33	Bangkok	5.782	1569.00	3.68514978
## 34	Seoul	9.806	605.00	16.20826446
## 35	Nagoya	2.320	326.00	7.11656442
## 36	Hyderabad	6.993	650.00	10.75846154
## 37	London	8.825	1572.00	5.61386768
## 38	Tehran	9.033	751.00	12.02796272
## 39	Chicago	2.706	589.00	4.59422750
## 40	Chengdu	16.045	14378.00	1.11594102
## 41	Nanjing	7.260	6582.00	1.10300820
## 42	Wuhan	10.893	8494.00	1.28243466
## 43	Ho Chi Minh City	7.431	2061.00	3.60553130
## 44	Luanda	2.166	116.00	18.67241379
## 45	Ahmedabad	5.571	464.00	12.00646552
## 46	Kuala Lumpur	1.768	243.00	7.27572016
## 47	Xi'an	8.989	10135.00	0.88692649
## 48	Hong Kong	7.299	1104.00	6.61141304
## 49	Dongguan	8.342	2465.00	3.38417850
## 50	Hangzhou	9.468	16596.00	0.57049892
## 51	Foshan	7.197	3848.00	1.87032225
## 52	Shenyang	8.294	1298.00	6.38983051
## 53	Riyadh	6.694	1913.00	3.49921589
## 54	Baghdad	8.127	5200.00	1.56288462
## 55	Santiago	0.236	22.00	10.72727273
## 56	Surat	4.467	327.00	13.66055046
## 57	Madrid	3.266	606.00	5.38943894
## 58	Suzhou	10.722	8488.42	1.26313260
## 59	Pune	3.124	276.00	11.31884058
## 60	Harbin	10.636	53068.00	0.20042210
## 61	Houston	2.326	1553.00	1.49774630
## 62	Dallas	1.345	882.00	1.52494331
## 63	Toronto	2.732	630.00	4.33650794
## 64	Dar es Salaam	4.365	1393.00	3.13352477
## 65	Miami	0.471	92.90	5.06996771
## 66	Belo Horizonte	2.503	330.90	7.56421880
## 67	Singapore	5.639	725.70	7.77042855
## 68	Philadelphia	1.526	369.59	4.12889959
## 69	Atlanta	0.420	354.22	1.18570380
## 70	Fukuoka	1.589	343.39	4.62739160
## 71	Khartoum	0.640	22142.00	0.02890434
## 72	Barcelona	1.620	101.40	15.97633136
## 73	Johannesburg	NA	NA	NA
## 74	Saint Petersburg	NA	NA	NA
## 75	Qingdao	NA	NA	NA
## 76	Dalian	NA	NA	NA
## 77	Washington, D.C.	0.702	177.00	3.96610169

```
## 78      Yangon      NA      NA      NA
## 79    Alexandria      NA      NA      NA
## 80      Jinan    8.700 10244.00 0.84927763
## 81    Guadalajara  1.460 10244.00 0.14252245
```

```
# d. Calculate the average city density by continent
```

```
average_density <- cities %>%
```

```
  # Group by continent
```

```
  group_by(continent) %>%
```

```
  # Calculate average, handling NA values
```

```
  summarise(average_city_density = mean(city_density, na.rm = TRUE))
```

```
# Print the result
```

```
print(average_density)
```

```
## # A tibble: 5 x 2
```

```
##   continent      average_city_density
```

```
##   <chr>          <dbl>
```

```
## 1 Africa              5.21
```

```
## 2 Asia              10.6
```

```
## 3 Europe             9.26
```

```
## 4 North America      3.87
```

```
## 5 South America      7.87
```

```
# e. Calculate metro_density if not already available
```

```
cities <- cities %>%
```

```
  mutate(metro_density = metro_pop / metro_area)
```

```
# Filter out rows with missing or zero values for city_density or metro_density to avoid issues with log scale
```

```
filtered_cities <- cities %>%
```

```
  filter(!is.na(city_density), !is.na(metro_density), city_density > 0, metro_density > 0)
```

```
# Create the plot
```

```
ggplot(filtered_cities, aes(x = city_density, y = metro_density, label = name)) +
```

```
  geom_point(alpha = 0.7) +
```

```
  geom_text_repel(size = 3, alpha = 0.7) +
```

```
  scale_x_log10() +
```

```
  scale_y_log10() +
```

```
  theme_minimal() +
```


```
  labs(title = "City Density vs Metro Density (Log Scale)",
```

```
        x = "City Density (per 1000 sq km, log scale)",
```

```
        y = "Metro Density (per sq km, log scale)")
```

```
## Warning: ggrepel: 2 unlabeled data points (too many overlaps). Consider
```

```
## increasing max.overlaps
```




Question 4

```
# Load the dataset
economist_data <- read.csv("EconomistData.csv")

# a. Create a scatter plot with the appropriate columns
ggplot(economist_data, aes(x = Percent.of.15plus.with.bank.account, y = SEDA.Current.level)) +
  geom_point() +
  theme_minimal() +
  labs(title = "Scatter Plot of Bank Account Ownership vs SEDA Score",
       x = "Percent of People (15+) with Bank Account",
       y = "SEDA Current Level")
```

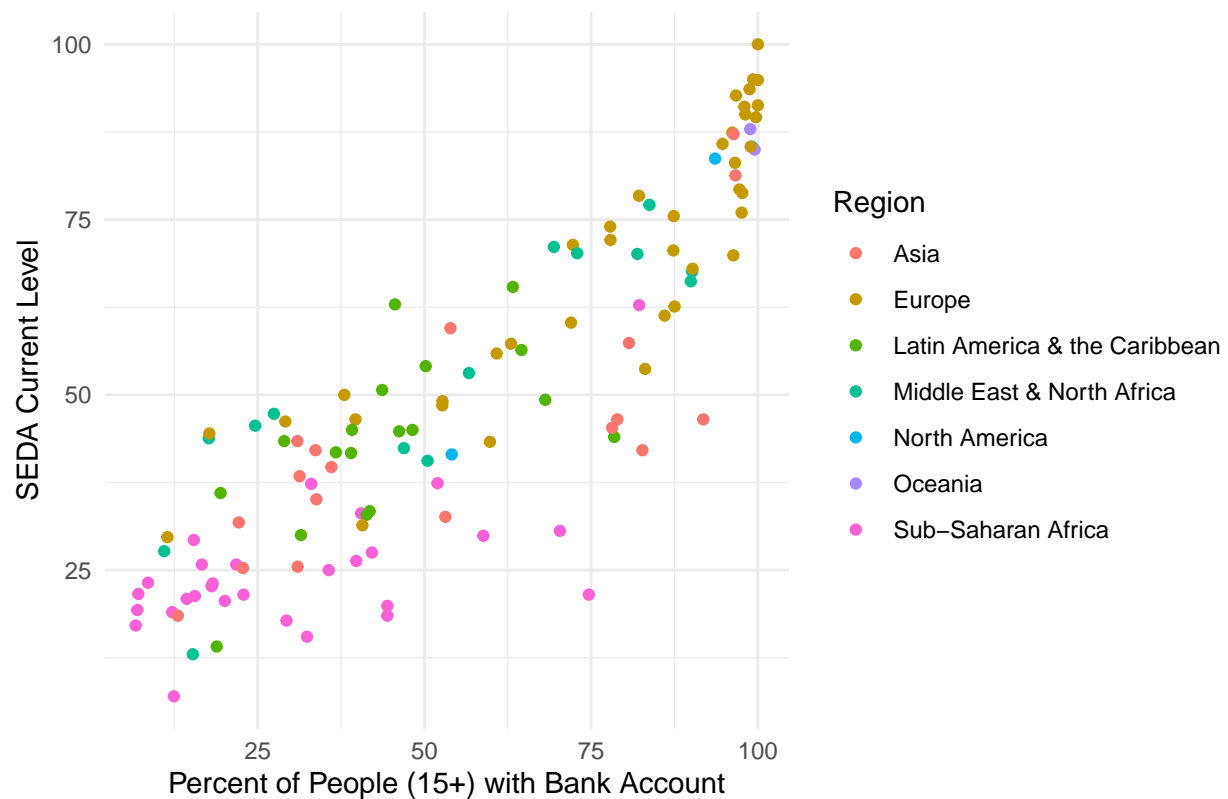



```
# b. Color all points blue.
ggplot(economist_data, aes(x = Percent.of.15plus.with.bank.account, y = SEDA.Current.level)) +
  geom_point(color = "blue") +
  theme_minimal() +
  labs(title = "Scatter Plot of Bank Account Ownership vs SEDA Score",
       x = "Percent of People (15+) with Bank Account",
       y = "SEDA Current Level")
```



```
# c. Color points according to the Region variable.
ggplot(economist_data, aes(x = Percent.of.15plus.with.bank.account, y = SEDA.Current.level, color = Reg
  geom_point() +
  theme_minimal() +
  labs(title = "Scatter Plot of Bank Account Ownership vs SEDA Score",
        x = "Percent of People (15+) with Bank Account",
        y = "SEDA Current Level")
```

Scatter Plot of Bank Account Ownership vs SEDA Score



```
# d. Overlay a fitted smoothing trend on top of the scatter plot.
ggplot(economist_data, aes(x = Percent.of.15plus.with.bank.account, y = SEDA.Current.level, color = Region)) +
  geom_point() +
  # Adding a fitted smoothing trend with low span
  geom_smooth(span = 0.3, method = "loess", se = FALSE) +
  theme_minimal() +
  labs(title = "Scatter Plot with Fitted Smoothing Trend (Low Span)",
       x = "Percent of People (15+) with Bank Account",
       y = "SEDA Current Level")
```


```
## 'geom_smooth()' using formula = 'y ~ x'
```

```
## Warning in sqrt(sum.squares/one.delta): NaNs produced
```

```
## Warning: Failed to fit group 5.
## Caused by error in 'simpleLoess()':
## ! span is too small
```

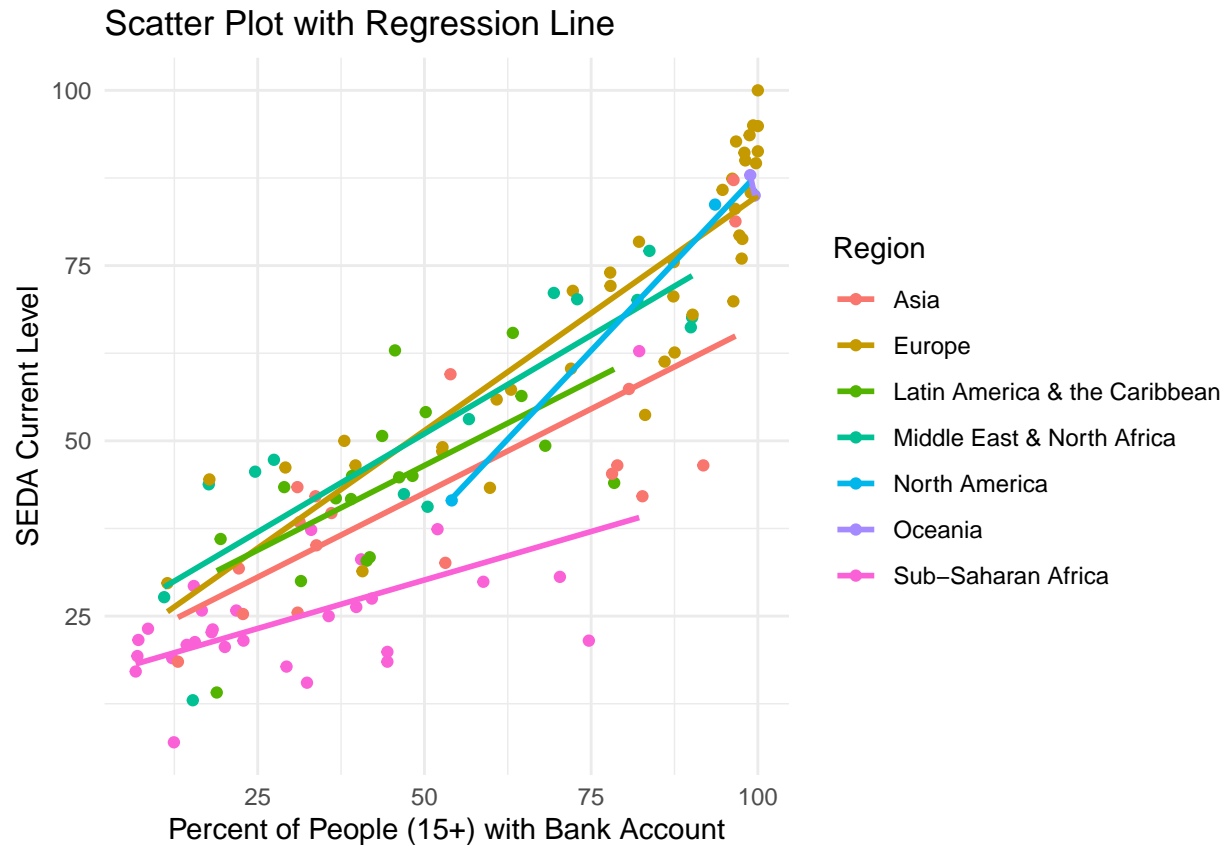
```
## Warning: Failed to fit group 6.
## Caused by error in 'simpleLoess()':
## ! span is too small
```

Scatter Plot with Fitted Smoothing Trend (Low Span)



```
# e. Create a scatter plot with a regression line overlaid
ggplot(economist_data, aes(x = Percent.of.15plus.with.bank.account, y = SEDA.Current.level, color = Region)) +
  geom_point() +
  # Overlay a regression line (method = "lm" for linear regression)
  geom_smooth(method = "lm", se = FALSE) +
  theme_minimal() +
  labs(title = "Scatter Plot with Regression Line",
        x = "Percent of People (15+) with Bank Account",
        y = "SEDA Current Level")
```


```
## 'geom_smooth()' using formula = 'y ~ x'
```



```
# f. Create a scatter plot with a regression line and facet it by Region
ggplot(economist_data, aes(x = Percent.of.15plus.with.bank.account, y = SEDA.Current.level, color = Region)) +
  geom_point() +
  # Overlay a regression line (method = "lm" for linear regression)
  geom_smooth(method = "lm", se = FALSE) +
  # Facet the plot by Region
  facet_wrap(~ Region) +
  theme_minimal() +
  labs(title = "Scatter Plot of Bank Account Ownership vs SEDA Score Faceted by Region",
       x = "Percent of People (15+) with Bank Account",
       y = "SEDA Current Level")
```

```
## 'geom_smooth()' using formula = 'y ~ x'
```

Scatter Plot of Bank Account Ownership vs SEDA Score Faceted by Region




Question 5

```
# Load the ggplot2 library
library(ggplot2)
# Load the mtcars dataset
data(mtcars)

# a. Convert the `hp` (horsepower) variable into a factor with three levels: "Low", "Medium", "High"
mtcars$hp_level <- cut(mtcars$hp,
                       breaks = c(-Inf, 100, 200, Inf),
                       labels = c("Low", "Medium", "High"))

# b. Create the scatter plot of `mpg` vs `wt`, faceted by the new `hp_level` factor
ggplot(mtcars, aes(x = wt, y = mpg)) +
  geom_point() +
  facet_wrap(~ hp_level) +
  theme_minimal() +
  labs(title = "Scatter Plot of MPG vs Weight, Faceted by Horsepower Level",
       x = "Weight (1000 lbs)",
       y = "Miles per Gallon (mpg)")
```

Scatter Plot of MPG vs Weight, Faceted by Horsepower Level



c. How does converting hp into categorical groups enhance the interpretability of the plot?

Converting hp into categories (Low, Medium, High) allows for easier comparison between groups rather than interpreting individual values across a continuous scale.

d. Describe the differences observed in mpg for different hp levels.

Vehicles in the "Low" horsepower category tend to have higher mpg values, indicating they are more fuel-efficient.

Vehicles with "Medium" horsepower have a moderate mpg value, showing a balance between power and efficiency

Vehicles in the "High" horsepower group generally have lower mpg, indicating less fuel efficiency due to their higher power requirements.

e. What function is used to create categorical levels from continuous variables?

The cut() function is used to create categorical levels from continuous variables. It allows you to specify breakpoints and labels, making it easy to transform numeric data into discrete categories.

f. Can faceting by grouped levels provide more insight than using hp as a continuous variable on the x-axis?

It reduces the complexity by breaking the data into simpler, comparable segments.

```
# helps in focusing on trends within specific categories rather than analyzing each
# individual data point along a continuous scale.
```

Question 6

```
# Load libraries
library(dplyr)

# Load the dataset from the given URL
url <- "https://raw.githubusercontent.com/Juanets/movie-stats/master/movies.csv"
movies <- read.csv(url)

# a. Find a subset of the movies produced after 2005
movies.sub <- movies %>%
  filter(year > 2005)

# Display first few rows of the subset
head(movies.sub)
```

```
##              name rating   genre year
## 1      The Departed      R    Crime 2006
## 2  The Fast and the Furious: Tokyo Drift PG-13  Action 2006
## 3 Talladega Nights: the Ballad of Ricky Bobby PG-13  Comedy 2006
## 4      The Prestige PG-13    Drama 2006
## 5      Cars           G Animation 2006
## 6      300           R    Action 2006
##              released score  votes   director
## 1  October 6, 2006 (United States)  8.5 1200000  Martin Scorsese
## 2   June 16, 2006 (United States)  6.0  252000   Justin Lin
## 3   August 4, 2006 (United States)  6.6  172000   Adam McKay
## 4 October 20, 2006 (United States)  8.5 1200000 Christopher Nolan
## 5   June 9, 2006 (United States)  7.1  381000   John Lasseter
## 6   March 9, 2007 (United States)  7.6  750000   Zack Snyder
##              writer      star   country  budget   gross
## 1 William Monahan Leonardo DiCaprio United States 9.00e+07 291465373
## 2   Chris Morgan   Lucas Black United States 8.50e+07 158964610
## 3   Will Ferrell   Will Ferrell United States 7.25e+07 163362095
## 4 Jonathan Nolan   Christian Bale United Kingdom 4.00e+07 109676311
## 5   John Lasseter   Owen Wilson United States 1.20e+08 461991867
## 6   Zack Snyder    Gerard Butler United States 6.50e+07 456068181
##              company runtime
## 1      Warner Bros.    151
## 2   Universal Pictures    104
## 3   Columbia Pictures    108
## 4   Touchstone Pictures    130
## 5 Pixar Animation Studios    117
## 6      Warner Bros.    117
```

```
# b. Keep only the specified columns in `movies.sub`
movies.sub <- movies.sub %>%
  select(name, director, year, country, genre, budget, gross, score)
```



```
# Display first few rows of the modified subset
head(movies.sub)
```

```
##              name            director year
## 1      The Departed      Martin Scorsese 2006
## 2    The Fast and the Furious: Tokyo Drift      Justin Lin 2006
## 3 Talladega Nights: the Ballad of Ricky Bobby      Adam McKay 2006
## 4      The Prestige      Christopher Nolan 2006
## 5      Cars      John Lasseter 2006
## 6      300      Zack Snyder 2006
##      country    genre    budget    gross score
## 1  United States    Crime 9.00e+07 291465373 8.5
## 2  United States    Action 8.50e+07 158964610 6.0
## 3  United States    Comedy 7.25e+07 163362095 6.6
## 4  United Kingdom    Drama 4.00e+07 109676311 8.5
## 5  United States    Animation 1.20e+08 461991867 7.1
## 6  United States    Action 6.50e+07 456068181 7.6
```

```
# c. Calculate the profit for each movie in `movies.sub` as a fraction of its budget
movies.sub <- movies.sub %>%
```

```
  mutate(
    # Convert `budget` to million dollars and round to 1 decimal place
    budget = round(budget / 1e6, 1),
    # Convert `gross` to million dollars and round to 1 decimal place
    gross = round(gross / 1e6, 1),
    # Calculate profit as a fraction of the budget
    profit_fraction = (gross - budget) / budget
  )
```

```
# Display first few rows with profit calculation
head(movies.sub)
```

```
##              name            director year
## 1      The Departed      Martin Scorsese 2006
## 2    The Fast and the Furious: Tokyo Drift      Justin Lin 2006
## 3 Talladega Nights: the Ballad of Ricky Bobby      Adam McKay 2006
## 4      The Prestige      Christopher Nolan 2006
## 5      Cars      John Lasseter 2006
## 6      300      Zack Snyder 2006
##      country    genre    budget    gross score profit_fraction
## 1  United States    Crime    90.0  291.5    8.5      2.2388889
## 2  United States    Action   85.0  159.0    6.0      0.8705882
## 3  United States    Comedy   72.5  163.4    6.6      1.2537931
## 4  United Kingdom    Drama   40.0  109.7    8.5      1.7425000
## 5  United States    Animation 120.0  462.0    7.1      2.8500000
## 6  United States    Action   65.0  456.1    7.6      6.0169231
```

```
# d. Count the number of movies produced by each genre and order them in descending order
genre_count <- movies.sub %>%
```

```
  group_by(genre) %>%
  summarise(count = n()) %>%
  arrange(desc(count))
```

```
# Display the genre count
print(genre_count)
```

```
## # A tibble: 16 x 2
##   genre      count
##   <chr>    <int>
## 1 Action      738
## 2 Comedy      629
## 3 Drama       548
## 4 Biography   228
## 5 Animation   189
## 6 Crime       176
## 7 Adventure   151
## 8 Horror      136
## 9 Fantasy     10
## 10 Mystery     5
## 11 Sci-Fi      4
## 12 Thriller    4
## 13 Family      2
## 14 Musical     2
## 15 Romance     2
## 16 Sport       1
```

```
# e. Group the 'movies.sub' dataset by 'country' and 'genre', then calculate the required metrics
```

```
movies_grouped <- movies.sub %>%
  group_by(country, genre) %>%
  summarise(
    # Count the number of movies in each group
    count = n(),
    # Median of profit as a fraction of budget
    median_profit_fraction = median(profit_fraction, na.rm = TRUE),
    # Mean of the movie score
    mean_score = mean(score, na.rm = TRUE),
    # Variance of the movie score
    variance_score = var(score, na.rm = TRUE),
    # Drop the grouping after summarizing
    .groups = "drop"
  ) %>%
  # Arrange by count in descending order
  arrange(desc(count))

# Display the result
print(movies_grouped)
```

```
## # A tibble: 175 x 6
##   country      genre count median_profit_fraction mean_score variance_score
##   <chr>        <chr> <int>          <dbl>          <dbl>          <dbl>
## 1 United States Action    527          1.32          6.33          0.769
## 2 United States Comedy    502          1.42          6.14          0.798
## 3 United States Drama     306          1.66          6.56          0.757
## 4 United States Animat~  133          2.51          6.66          0.847
## 5 United States Biogra~  122          0.892          6.99          0.356
```

```
## 6 United States Crime 111 0.688 6.58 0.700
## 7 United States Horror 109 3.76 5.70 0.807
## 8 United States Advent~ 98 1.54 6.36 1.14
## 9 United Kingdom Drama 75 0.829 6.79 0.505
## 10 United Kingdom Action 67 0.547 6.65 0.560
## # i 165 more rows
```

Question 7

```
# library(dplyr)
movies_filtered <- movies %>%
  # Filter movies produced after 2001
  filter(year > 2001) %>%
  # Calculate the number of movies for each director and
  # filter out those with fewer than 4 movies
  group_by(director) %>%
  filter(n() >= 4) %>%
  ungroup() %>%
  # Group by genre and director, then calculate the mean score for each director
  group_by(genre, director) %>%
  summarise(mean_score = mean(score, na.rm = TRUE), .groups = "drop") %>%
  ungroup() %>%
  # Step 4: For each genre, select the top two directors with the highest mean scores
  group_by(genre) %>%
  top_n(n = 2, wt = mean_score) %>%
  # Arrange by genre and descending mean score for clarity
  arrange(genre, desc(mean_score)) %>%
  ungroup()
  .groups = "drop"

# Display the result
print(movies_filtered)
```

```
## # A tibble: 25 x 3
##   genre      director      mean_score
##   <chr>      <chr>          <dbl>
## 1 Action    Park Chan-Wook      8.4
## 2 Action    Christopher Nolan    8.27
## 3 Adventure Christopher Nolan    8.6
## 4 Adventure Quentin Tarantino 8.3
## 5 Animation Andrew Stanton    7.93
## 6 Animation Wes Anderson    7.9
## 7 Biography Roman Polanski    8.5
## 8 Biography Tom McCarthy    8.1
## 9 Comedy    Bong Joon Ho        8.6
## 10 Comedy    Rian Johnson        7.9
## # i 15 more rows
```

Question 8

The continuous random variable X has the following probability density function (PDF), for some positive constant c :

$$f(x) = \frac{3}{(1+x)^3}, \quad 0 \leq x \leq c$$

To determine the constant c such that $f(x)$ is a legitimate PDF, we need to ensure that the total area under the curve of $f(x)$ over the given domain is equal to 1.

Step 1: Set Up the Integral

To find the value of c :

$$\int_0^c f(x) dx = 1$$

Substituting $f(x)$:

$$\int_0^c \frac{3}{(1+x)^3} dx = 1$$

Step 2: Solve the Integral

To evaluate this integral, we use substitution. Let:

$$u = 1 + x, \quad \text{then } du = dx$$

Rewrite the limits in terms of u :

- When $x = 0$, $u = 1$
- When $x = c$, $u = 1 + c$

The integral becomes:

$$\int_1^{1+c} \frac{3}{u^3} du$$

Now, evaluate the antiderivative:

$$\int \frac{3}{u^3} du = 3 \int u^{-3} du = 3 \cdot \left(\frac{u^{-2}}{-2} \right) = -\frac{3}{2} u^{-2}$$

Step 3: Substitute the Limits

Evaluate the definite integral from $u = 1$ to $u = 1 + c$:

$$\left[-\frac{3}{2} \cdot \frac{1}{u^2} \right]_1^{1+c} = -\frac{3}{2(1+c)^2} + \frac{3}{2}$$

Simplify:

$$\frac{3}{2} - \frac{3}{2(1+c)^2} = 1$$

Step 4: Solve for c

To find c :

1. Move the constant to the other side:

$$\frac{3}{2} - 1 = \frac{3}{2(1+c)^2}$$

$$\frac{1}{2} = \frac{3}{2(1+c)^2}$$

2. Cross multiply to solve for $(1+c)^2$:

$$1 \cdot 2(1+c)^2 = 2 \cdot 3$$

$$2(1+c)^2 = 6$$

3. Divide both sides by 2:

$$(1+c)^2 = 3$$

4. Take the square root of both sides:

$$1+c = \sqrt{3}$$

5. Solve for c :

$$c = \sqrt{3} - 1$$

Conclusion

The value of c that makes $f(x)$ a legitimate PDF is:

$$c = \sqrt{3} - 1$$

Now that we have the value of c , we can plot the PDF in R:

```

# (b)
# Define the value of c
c <- sqrt(3) - 1


# Define the PDF function
f <- function(x) {
  3 / (1 + x)^3
}

# Define the range for plotting
x_vals <- seq(0, c, length.out = 1000)

# Calculate the y-values for each x
y_vals <- f(x_vals)

# Plot the PDF
plot(x_vals, y_vals, type = "l", col = "blue", lwd = 2,
     main = "Probability Density Function of X",
     xlab = "x", ylab = "f(x)")

```



(c) The Expected Value $E(X)$

The expected value $E(X)$ is defined as:

$$E(X) = \int_0^c x \cdot f(x) dx = \int_0^c x \cdot \frac{3}{(1+x)^3} dx$$

Step 2.1: Set Up and Simplify the Integral

Substituting the value of $f(x)$:

$$E(X) = \int_0^c \frac{3x}{(1+x)^3} dx$$

We use substitution again:

- Let $u = 1 + x$, then $du = dx$, and $x = u - 1$.
- When $x = 0$, $u = 1$.
- When $x = c$, $u = 1 + c$.

The integral becomes:

$$E(X) = \int_1^{1+c} \frac{3(u-1)}{u^3} du = 3 \int_1^{1+c} \left(\frac{u-1}{u^3} \right) du$$

Expanding:

$$E(X) = 3 \int_1^{1+c} \left(\frac{1}{u^2} - \frac{1}{u^3} \right) du$$

Step 2.2: Integrate Each Term

Integrate each term:

$$\int u^{-2} du = -\frac{1}{u}$$

$$\int u^{-3} du = -\frac{1}{2u^2}$$

Evaluating the definite integral:

$$E(X) = 3 \left(\left[-\frac{1}{u} \right]_1^{1+c} - \left[-\frac{1}{2u^2} \right]_1^{1+c} \right)$$

Substituting the limits:

1. **For** $-\frac{1}{u}$:

$$\left[-\frac{1}{u} \right]_1^{1+c} = -\frac{1}{1+c} + \frac{1}{1} = 1 - \frac{1}{1+c}$$

2. **For** $-\frac{1}{2u^2}$:

$$\left[-\frac{1}{2u^2}\right]_1^{1+c} = -\frac{1}{2(1+c)^2} + \frac{1}{2}$$

Putting it all together:

$$E(X) = 3 \left(1 - \frac{1}{1+c} + \frac{1}{2} - \frac{1}{2(1+c)^2} \right)$$

Step 2.3: Substitute $c = \sqrt{3} - 1$

Substituting $c = \sqrt{3} - 1$:

1. $1 + c = \sqrt{3}$

$$E(X) = 3 \left(1 + \frac{1}{2} - \frac{1}{\sqrt{3}} - \frac{1}{2(\sqrt{3})^2} \right)$$

Simplify:

- $1 + \frac{1}{2} = \frac{3}{2}$
- $\frac{1}{2(\sqrt{3})^2} = \frac{1}{6}$

Combining:

$$E(X) = 3 \left(\frac{3}{2} - \frac{1}{\sqrt{3}} - \frac{1}{6} \right)$$

Common denominator for $\frac{3}{2}$ and $\frac{1}{6}$:

$$\frac{3}{2} - \frac{1}{6} = \frac{8}{6} = \frac{4}{3}$$

So:

$$E(X) = 3 \left(\frac{4}{3} - \frac{1}{\sqrt{3}} \right)$$

Multiply by 3:

$$E(X) = 4 - \frac{3}{\sqrt{3}}$$

Simplify:

$$E(X) = 4 - \sqrt{3}$$

Final Answer

The expected value $E(X)$ is:

$$E(X) = 4 - \sqrt{3}$$

$$E(X) = 2.26$$

#(d) Define the function for the cumulative distribution (inverse transformation method)

```
func_inv <- function(u) {  
  (2 * (1 - u))(-1/2) - 1  
}
```

Simulate 1000 random observations

```
set.seed(123) # Set a seed for reproducibility
```

```
u_vals <- runif(1000)
```

```
simulated_data <- func_inv(u_vals)
```

(e) Estimate mean and variance

```
estimated_mean <- mean(simulated_data)
```

```
estimated_variance <- var(simulated_data)
```

Print the results

```
cat("Estimated Mean:", estimated_mean, "\n")
```

```
## Estimated Mean: 0.3725972
```

```
cat("Estimated Variance:", estimated_variance, "\n")
```

```
## Estimated Variance: 1.944269
```